

CENG 331

Computer Organization

Fall '2020-2021

THE2: Attack Lab

Deadline: 06 December 2020, Sunday, 23:59

1 Introduction

This assignment involves performing a total of four attacks on two programs having different security vulnerabilities. Outcomes you will gain from this lab include:

- You will learn different ways that attackers can exploit security vulnerabilities when programs do not safeguard themselves well enough against buffer overflows.
- Through this, you will get a better understanding of how to write programs that are more secure, as well as some of the features provided by compilers and operating systems to make programs less vulnerable.
- You will gain a deeper understanding of the stack and parameter-passing mechanisms of x86-64 machine code.
- You will gain a deeper understanding of how x86-64 instructions are encoded.
- You will gain more experience with debugging tools such as `gdb` and `objdump`.

Note: In this lab, you will gain firsthand experience with methods used to exploit security weaknesses in operating systems and network servers. Our purpose is to help you learn about the runtime operation of programs and to understand the nature of these security weaknesses so that you can avoid them when you write system code. We do not condone the use of any other form of attack to gain unauthorized access to any system resources.

2 Specifications

As usual, this is an **individual** project. You will generate attacks for target programs that are custom generated for you.

2.1 Getting Files

You can obtain your files by pointing your Web browser at:

`http://144.122.71.101:8083/`

The server will build your files and return them to your browser in a **tar** file called **target k .tar**, where k is the unique number of your target programs.

Note: It takes a few seconds to build and download your target, so please be patient.

Save the **target k .tar** file in a (protected) Linux directory in which you plan to do your work. Then give the command: **tar -xvf target k .tar**. This will extract a directory **target k** containing the files described below.

You should only download one set of files. If for some reason you download multiple targets, choose one target to work on and delete the rest.

The files in **target k** include:

README.txt: A file describing the contents of the directory

ctarget: An executable program vulnerable to *code-injection* attacks

rtarget: An executable program vulnerable to *return-oriented-programming* attacks

cookie.txt: An 8-digit hex code that you will use as a unique identifier in your attacks.

farm.c: The source code of your target's "gadget farm", which you will use in generating return-oriented programming attacks.

hex2raw: A utility to generate attack strings.

In the following instructions, we will assume that you have copied the files to a protected local directory, and that you are executing the programs in that local directory.

2.2 Grading

Similar to the first assignment, Attack lab will consist of two parts. The first part is this homework and the second part is a quiz which will test your familiarity with the concepts and tools used in the homework. The quiz will be graded over 10 points and your final grade will be calculated by scaling the homework grade with the quiz grade.

$$\text{AttackLabGrade} = (\text{HomeworkGrade} / 10) * (\text{QuizGrade} + 5)$$

The quiz will take place on 7 December, Monday, 18:00

2.3 Important Points

Here is a summary of some important rules regarding valid solutions for this lab. These points will not make much sense when you read this document for the first time. They are presented here as a central reference of rules once you get started.

- You must do the assignment on a machine that is similar to the one that generated your targets.
- Your solutions may not use attacks to circumvent the validation code in the programs. Specifically, any address you incorporate into an attack string for use by a `ret` instruction should be to one of the following destinations:
 - The addresses for functions `touch1`, `touch2`, or `touch3`. For `touch1` you are allowed to select an address inside the function.
 - The address of your injected code
 - The address of one of your gadgets from the gadget farm.
- You may only construct gadgets from file `rtarget` with addresses ranging between those for functions `start_farm` and `end_farm`.

3 Target Programs

Both CTARGET and RTARGET read strings from standard input. They do so with the function `getbuf` defined below:

```
1 unsigned getbuf() {
2     char buf[BUFFER_SIZE];
3     Gets(buf);
4     return 1;
5 }
```

The function `Gets` is similar to the standard library function `gets`. It reads a string from standard input terminated by EOF and stores it (along with a null terminator) at the specified destination. Unlike the standard library function it doesn't stop at the newline character. In this code, you can see that the destination is an array `buf`, declared as having `BUFFER_SIZE` bytes. At the time your targets were generated, `BUFFER_SIZE` was a compile-time constant specific to your version of the programs.

Functions `Gets()` and `gets()` have no way to determine whether their destination buffers are large enough to store the string they read. They simply copy sequences of bytes, possibly overrunning the bounds of the storage allocated at the destinations.

If the string typed by the user and read by `getbuf` is sufficiently short, it is clear that `getbuf` will return 1, as shown by the following execution examples:

```
unix> ./ctarget
Cookie: 0x1a7dd803
Type string: Keep it short!
[enter CTRL+D after newline, it will terminate here]
No exploit.
Getbuf returned 0x1 Normal return
```

Typically an error occurs if you type a long string:

```
unix> ./ctarget
Cookie: 0x1a7dd803
Type string: This is not a very interesting string, but it has the property ...
[enter CTRL+D after newline, it will terminate here]
Ouch!: You caused a segmentation fault!
Better luck next time
```

(Note that the value of the cookie shown will differ from yours.) Program RTARGET will have the same behavior. As the error message indicates, overrunning the buffer typically causes the program state to be corrupted, leading to a memory access error. Your task is to be more clever with the strings you feed CTARGET and RTARGET so that they do more interesting things. These are called *exploit* strings.

Both CTARGET and RTARGET take several different command line arguments:

- h: Print list of possible command line arguments
- q: Don't send results to the grading server. Offline working option.
- i FILE: Supply input from a file, rather than from standard input

You can use -q option for offline checking of your result. You can also use this command in gdb when running your code. To run your code offline, you can give the -q parameter with run command in gdb. You can similarly run your code with -i parameter in gdb.

Example:

```
unix> gdb ./ctarget
(gdb) r -q
(gdb) r -i ctarget.l1.raw
(gdb) r -q -i ctarget.l1.raw
```

Your exploit strings will typically contain byte values that do not correspond to the ASCII values for printing characters. The program HEX2RAW will enable you to generate these *raw* strings. See Appendix A for more information on how to use HEX2RAW.

Important points:

- The **Gets** function will only stop when it encounters an EOF, and since this is not a character that has an ASCII value, your exploit string cannot be cut short because of any character. However this means that when you are testing your target program by hand, you need to terminate it with an EOF which you can send it using CTRL+D key combination. Example:

```
unix> ./ctarget
Cookie: 0x1a7dd803
Type string: exploit string
[enter CTRL+D after newline, it will terminate here]
No exploit. Getbuf returned 0x1
Normal return
```

- HEX2RAW expects two-digit hex values separated by one or more white spaces. So if you want to create a byte with a hex value of 0, you need to write it as 00. To create the word 0xdeadbeef you should pass "ef be ad de" to HEX2RAW. Note the reversal required for little-endian byte ordering.

When you have correctly solved one of the levels, your target program will automatically send a notification to the grading server. For example:

```
unix> ./hex2raw < ctargget.l2 | ./ctargget
Cookie: 0x1a7dd803
Type string:Touch2!: You called touch2(0x1a7dd803, 0x2b5cd803)
Valid solution for level 2 with target ctargget
PASSED: Sent exploit string to server to be validated.
NICE JOB!
```

The server will test your exploit string to make sure it really works, and it will update the Attacklab scoreboard page indicating that your `userid` (listed by your target number for anonymity) has completed this phase.

Important Note: Unfortunately there is a problem with `ctargget` that results in a segmentation fault after `Valid solution for level X with target ctargget` text. If you see your grade in the scoreboard, then your solution is valid. You can ignore the error message. For example:

```
unix> ./hex2raw < ctargget.l2 | ./ctargget
Cookie: 0x1a7dd803
Type string:Touch2!: You called touch2(0x1a7dd803, 0x2b5cd803)
Valid solution for level 2 with target ctargget
Segmentation Fault
FAILED!
```

This problem occurs because of a `movaps` instruction inside the validation part of the program. This instruction expects the data on the stack to be 16-bit aligned. In other words, it expects the last digit of `%rsp` to be 0. However, since we alter the stack and `%rsp`, we fail to satisfy this requirement for `movaps` instruction and cause a segmentation fault. There is a way to construct an exploit string that won't cause a segmentation fault, but you are not expected to craft your exploit strings in that way.

You can view the scoreboard by pointing your Web browser at

<http://144.122.71.101:8083/scoreboard>

Unlike the Bomb Lab, there is no penalty for making mistakes in this lab. Feel free to fire away at CTARGET and RTARGET with any strings you like. You can also find your solutions in your own linux machines with offline mode and use inek machines to send your final solutions after.

Level	Program	Method	Function	Points
1	CTARGET	CI	<code>touch1</code>	10
2	CTARGET	CI	<code>touch2</code>	25
3	CTARGET	CI	<code>touch3</code>	30
4	RTARGET	ROP	<code>touch2</code>	35

CI: Code injection

ROP: Return-oriented programming

Figure 1: Summary of attack lab levels

Figure 1 summarizes the four levels of the lab. As can be seen, the first three involve code-injection (CI) attacks on CTARGET, while the last one involve return-oriented-programming (ROP) attack on RTARGET.

4 Part I: Code Injection Attacks

For the first three levels, your exploit strings will attack CTARGET. This program is set up in a way that the stack positions will be consistent from one run to the next and so that data on the stack can be treated as executable code. These features make the program vulnerable to attacks where the exploit strings contain the byte encodings of executable code.

4.1 Level 1

For Level 1, you will not inject new code. Instead, your exploit string will redirect the program to execute an existing procedure.

Function `getbuf` is called within CTARGET by a function `test` having the following C code:

```
1 void test()
2 {
3     int val;
4     val = getbuf();
5     printf("No exploit.  Getbuf returned 0x%x\n", val);
6 }
```

When `getbuf` executes its return statement (line 4 of `getbuf`), the program ordinarily resumes execution within function `test` (at line 5 of this function). We want to change this behavior. Within the file `ctarget`, there is code for a function `touch1` having the following C representation:

```
1 void touch1() {
2     srand(0x1337);
3     if (rand() % 0xBADA55 == 0xB04FE4) { // Always returns 0xB04FE4.
4         printf(
5             "Misfire: You called touch1() but you must not execute this
6             part.\n");
7         fail(1);
8     } else {
9         if (rand() % 0xF0015 == 0x37513) { // Execute this part.
10            vlevel = 1;                      /* Part of validation protocol
11            */
12            printf("Touch1!: You called touch1() correctly.\n");
13            validate(1);
14        } else { // This will always execute.
15            vlevel = 1;
16            printf(
17                "Misfire: You called touch1() but you must not execute this
18                part.\n");
19            fail(1);
20        }
21    }
22    exit(0);
23 }
```

Your task is to get CTARGET to execute the code for `touch1` when `getbuf` executes its return statement, rather than returning to `test`. However you should make sure the if path that results in function call `validate(1)` is executed. Please note that `touch1` function will always execute in a way that will fail if not redirected. Note that your exploit string may also corrupt parts of the stack not directly related to this stage, but this will not cause a problem, since `touch1` causes the program to exit directly.

Some Advice:

- All the information you need to devise your exploit string for this level can be determined by examining a disassembled version of CTARGET. Use `objdump -d` to get this disassembled version.
- The idea is to position a byte representation of the address where correct `vlevel=1` instruction is executed so that the `ret` instruction at the end of the code for `getbuf` will transfer control to it.
- Be careful about byte ordering.
- You might want to use GDB to step the program through the last few instructions of `getbuf` to make sure it is doing the right thing.
- The placement of `buf` within the stack frame for `getbuf` depends on the value of compile-time constant `BUFFER_SIZE`, as well the allocation strategy used by GCC. You will need to examine the disassembled code to determine its position.

4.2 Level 2

Level 2 involves injecting a small amount of code as part of your exploit string.

Within the file `ctarget` there is code for a function `touch2` having the following C representation:

```
1 void touch2(unsigned v1, unsigned v2) {
2     vlevel = 2; /* Part of validation protocol */
3     if (v1 == cookie && v2 == (cookie & ((cookie >> 16) | 0xAD0BE000))) {
4         printf("Touch2!: You called touch2(0x%.8x, 0x%.8x)\n", v1, v2);
5         validate(2);
6     } else {
7         printf("Misfire: You called touch2(0x%.8x, 0x%.8x)\n", v1, v2);
8         fail(2);
9     }
10    exit(0);
11 }
```

Your task is to get CTARGET to execute the code for `touch2` rather than returning to `test`. In this case, however, you must make it appear to `touch2` as if you have passed your `cookie` as its first argument and `cookie & ((cookie >> 16) | 0xAD0BE000)` as your second argument. Do not worry about the overflow.

Some Advice:

- You will want to position a byte representation of the address of your injected code in such a way that `ret` instruction at the end of the code for `getbuf` will transfer control to it.
- Recall that the first argument to a function is passed in register `%rdi` and the second argument is passed in register `%rsi`.
- Your injected code should set the registers to your `cookie` and `cookie & ((cookie >> 16) | 0xADOBE000)`, then use a `ret` instruction to transfer control to the first instruction in `touch2`.
- Do not attempt to use `jmp` or `call` instructions in your exploit code. The encodings of destination addresses for these instructions are difficult to formulate. Use `ret` instructions for all transfers of control, even when you are not returning from a call.
- See the discussion in Appendix B on how to use tools to generate the byte-level representations of instruction sequences.

4.3 Level 3

Level 3 also involves a code injection attack, but passing a string as its first argument and unsigned int array of size 8 as its second argument where the pointer should point to its first element (its address).

Within the file `ctarget` there is code for functions `hexmatch`, `checkmult` and `touch3` having the following C representations:

C representations of `hexmatch`, `checkmult` and `touch3` are given below.

Your task is to get `CTARGET` to execute the code for `touch3` rather than returning to `test`. You must make it appear to `touch3` as if you have passed two arguments. First argument must be the string representation of your cookie. Second argument should be an unsigned array of size 8 where each value at index i is the multiplication of the ASCII values of your cookie's i and $i - 1$ indexes. For index $i = 0$, instead of $i - 1$, character at the last index of your cookie is used as the second operand of the multiplication operation

Some Advice:

- You will either need to include a string representation of your cookie in your exploit string or write an assembly code to put your representation in the stack. The string should consist of the eight hexadecimal digits (ordered from most to least significant) without a leading “0x.”
- Recall that a string is represented in C as a sequence of bytes followed by a byte with value 0. Type “`man ascii`” on any Linux machine to see the byte representations of the characters you need. We have also included an ASCII table in your references.
- Second argument should have a 8 unsigned int characters consecutively. Please also note that unsigned integers are 4 bytes long.
- Your injected code should set register `%rdi` to the address of this cookie string and `%rsi` to the address of the multiplication array.
- When functions `hexmatch`, `checkmult` and `strncmp` are called, they push data onto the stack, overwriting portions of memory that held the buffer used by `getbuf`. As a result, you will need to be careful where you place the string representation of your cookie and the array.


```

1 /* Compare string to hex representation of unsigned value */
2 int hexmatch(unsigned val, char *sval) {
3     if (sval == NULL) return 0;
4     char cbuf[120];
5     /* Make position of check string unpredictable */
6     char *s = cbuf + random() % 110;
7     sprintf(s, "%.8x", val);
8     return strncmp(sval, s, 9) == 0;
9 }
10 /* Check if the cookie and its consecutive values are multiplied
11     correctly in each element. The size of the mult array is 8. */
12 int checkmult(unsigned cookie_param, unsigned *cookie_consec_mult) {
13     if (cookie_consec_mult == NULL) return 0;
14     char cbuf[70];
15     /* Make position of check multiplication unpredictable */
16     char *s = cbuf + random() % 60;
17     sprintf(s, "%.8x", cookie_param);
18     for (int i = 0, j = 0; i < 8; i++) { // Array size is 6
19         j = i - 1;
20         if (i == 0) {
21             j = 7;
22         }
23         if (cookie_consec_mult[i] != s[i] * s[j])
24             return 0;
25     }
26     return 1;
27 }
28
29 void touch3(char *cookie_string, unsigned *cookie_consec_mult) {
30     vlevel = 3; /* Part of validation protocol */
31     if (hexmatch(cookie, cookie_string) &&
32         checkmult(cookie, cookie_consec_mult)) {
33         printf("Touch3!: You called touch3 with correct parameters.\n");
34         validate(3);
35     } else {
36         printf("Misfire: You called touch3 with %s as your first "
37             "argument and %p as the address of your second argument.\n",
38             cookie_string, (void *)cookie_consec_mult);
39         printf("Contents of your second argument follows below:\n");
40         fflush(stdout);
41         for (int i = 0; i < 6; i++) {
42             printf("%.8x ", cookie_consec_mult[i]);
43             fflush(stdout);
44         }
45         printf("\n");
46         fail(3);
47     }
48     exit(0);
49 }

```

5 Part II: Return-Oriented Programming

Performing code-injection attacks on program `RTARGET` is much more difficult than it is for `CTARGET`, because it uses two techniques to thwart such attacks:

- It uses randomization so that the stack positions differ from one run to another. This makes it impossible to determine where your injected code will be located.
- It marks the section of memory holding the stack as nonexecutable, so even if you could set the program counter to the start of your injected code, the program would fail with a segmentation fault.




Figure 2: Setting up sequence of gadgets for execution. Byte value `0xc3` encodes the `ret` instruction.

Fortunately, clever people have devised strategies for getting useful things done in a program by executing existing code, rather than injecting new code. The most general form of this is referred to as *return-oriented programming* (ROP). The strategy with ROP is to identify byte sequences within an existing program that consist of one or more instructions followed by the instruction `ret`. Such a segment is referred to as a *gadget*. Figure 2 illustrates how the stack can be set up to execute a sequence of n gadgets. In this figure, the stack contains a sequence of gadget addresses. Each gadget consists of a series of instruction bytes, with the final one being `0xc3`, encoding the `ret` instruction. When the program executes a `ret` instruction starting with this configuration, it will initiate a chain of gadget executions, with the `ret` instruction at the end of each gadget causing the program to jump to the beginning of the next.

A gadget can make use of code corresponding to assembly-language statements generated by the compiler, especially ones at the ends of functions. In practice, there may be some useful gadgets of this form, but not enough to implement many important operations. For example, it is highly unlikely that a compiled function would have `popq %rdi` as its last instruction before `ret`. Fortunately, with a byte-oriented instruction set, such as x86-64, a gadget can often be found by extracting patterns from other parts of the instruction byte sequence.

For example, one version of `rtarget` contains code generated for the following C function:

```
1 void setval_210(unsigned *p) {  
2     *p = 3347663060U;  
3 }
```

The chances of this function being useful for attacking a system seem pretty slim. But, the disassembled machine code for this function shows an interesting byte sequence:

```

1 /* $begin farm-eg-d */
2 0000000000400f15 <setval_210>:
3   400f15: c7 07 d4 48 89 c7      movl    $0xc78948d4, (%rdi)
4   400f1b: c3                      retq
5 /* $end farm-eg-d */

```

The byte sequence 48 89 c7 encodes the instruction `movq %rax, %rdi`. (See Figure 3A for the encodings of useful `movq` instructions.) This sequence is followed by byte value c3, which encodes the `ret` instruction. The function starts at address 0x400f15, and the sequence starts on the fourth byte of the function. Thus, this code contains a gadget, having a starting address of 0x400f18, that will copy the 64-bit value in register `%rax` to register `%rdi`.

Your code for `RTARGET` contains a number of functions similar to the `setval_210` function shown above in a region we refer to as the *gadget farm*. Your job will be to identify useful gadgets in the gadget farm and use these to perform attacks similar to those you did in Level 2.

Important: The gadget farm is demarcated by functions `start_farm` and `end_farm` in your copy of `rtarget`. Do not attempt to construct gadgets from other portions of the program code.

5.1 Level 4

For Level 4, you will repeat the attack of Level 2, but do so on program `RTARGET` using gadgets from your gadget farm. You can construct your solution using gadgets consisting of the following instruction types, and using only the first eight x86-64 registers (`%rax-%rdi`).

`movq` : The codes for these are shown in Figure 3A.

`popq` : The codes for these are shown in Figure 3B.

`addq` : The codes for these are shown in Figure 3D.

`ret` : This instruction is encoded by the single byte 0xc3.

`nop` : This instruction (pronounced “no op,” which is short for “no operation”) is encoded by the single byte 0x90. Its only effect is to cause the program counter to be incremented by 1.

Some Advice:

- All the gadgets you need can be found in the region of the code for `rtarget` demarcated by the functions `start_farm` and `end_farm`.
- When a gadget uses a `popq` instruction, it will pop data from the stack. As a result, your exploit string will contain a combination of gadget addresses and data.

6 Handin

As a precaution, you need to submit a single file `hw2.tar`. The file should contain `ctarget.11`, `ctarget.12`, `ctarget.13` and `rtarget.12` files in textual format. These files should be same files that you feed to the `hex2raw` program and should be human readable.

Ex: `ctarget.11` First character after dot is the letter l not the number 1.

A. Encodings of movq instructions

movq S, D

Source S	Destination D							
	%rax	%rcx	%rdx	%rbx	%rsp	%rbp	%rsi	%rdi
%rax	48 89 c0	48 89 c1	48 89 c2	48 89 c3	48 89 c4	48 89 c5	48 89 c6	48 89 c7
%rcx	48 89 c8	48 89 c9	48 89 ca	48 89 cb	48 89 cc	48 89 cd	48 89 ce	48 89 cf
%rdx	48 89 d0	48 89 d1	48 89 d2	48 89 d3	48 89 d4	48 89 d5	48 89 d6	48 89 d7
%rbx	48 89 d8	48 89 d9	48 89 da	48 89 db	48 89 dc	48 89 dd	48 89 de	48 89 df
%rsp	48 89 e0	48 89 e1	48 89 e2	48 89 e3	48 89 e4	48 89 e5	48 89 e6	48 89 e7
%rbp	48 89 e8	48 89 e9	48 89 ea	48 89 eb	48 89 ec	48 89 ed	48 89 ee	48 89 ef
%rsi	48 89 f0	48 89 f1	48 89 f2	48 89 f3	48 89 f4	48 89 f5	48 89 f6	48 89 f7
%rdi	48 89 f8	48 89 f9	48 89 fa	48 89 fb	48 89 fc	48 89 fd	48 89 fe	48 89 ff

B. Encodings of popq instructions

Operation	Register R							
	%rax	%rcx	%rdx	%rbx	%rsp	%rbp	%rsi	%rdi
popq R	58	59	5a	5b	5c	5d	5e	5f

C. Encodings of movl instructions

movl S, D

Source S	Destination D							
	%eax	%ecx	%edx	%ebx	%esp	%ebp	%esi	%edi
%eax	89 c0	89 c1	89 c2	89 c3	89 c4	89 c5	89 c6	89 c7
%ecx	89 c8	89 c9	89 ca	89 cb	89 cc	89 cd	89 ce	89 cf
%edx	89 d0	89 d1	89 d2	89 d3	89 d4	89 d5	89 d6	89 d7
%ebx	89 d8	89 d9	89 da	89 db	89 dc	89 dd	89 de	89 df
%esp	89 e0	89 e1	89 e2	89 e3	89 e4	89 e5	89 e6	89 e7
%ebp	89 e8	89 e9	89 ea	89 eb	89 ec	89 ed	89 ee	89 ef
%esi	89 f0	89 f1	89 f2	89 f3	89 f4	89 f5	89 f6	89 f7
%edi	89 f8	89 f9	89 fa	89 fb	89 fc	89 fd	89 fe	89 ff

D. Encodings of addq instructions

addq S, D

Source S	Destination D							
	%rax	%rcx	%rdx	%rbx	%rsp	%rbp	%rsi	%rdi
%rax	48 01 c0	48 01 c1	48 01 c2	48 01 c3	48 01 c4	48 01 c5	48 01 c6	48 01 c7
%rcx	48 01 c8	48 01 c9	48 01 ca	48 01 cb	48 01 cc	48 01 cd	48 01 ce	48 01 cf
%rdx	48 01 d0	48 01 d1	48 01 d2	48 01 d3	48 01 d4	48 01 d5	48 01 d6	48 01 d7
%rbx	48 01 d8	48 01 d9	48 01 da	48 01 db	48 01 dc	48 01 dd	48 01 de	48 01 df
%rsp	48 01 e0	48 01 e1	48 01 e2	48 01 e3	48 01 e4	48 01 e5	48 01 e6	48 01 e7
%rbp	48 01 e8	48 01 e9	48 01 ea	48 01 eb	48 01 ec	48 01 ed	48 01 ee	48 01 ef
%rsi	48 01 f0	48 01 f1	48 01 f2	48 01 f3	48 01 f4	48 01 f5	48 01 f6	48 01 f7
%rdi	48 01 f8	48 01 f9	48 01 fa	48 01 fb	48 01 fc	48 01 fd	48 01 fe	48 01 ff

Figure 3: Byte encodings of instructions. All values are shown in hexadecimal.

A Using HEX2RAW

HEX2RAW takes as input a *hex-formatted* string. In this format, each byte value is represented by two hex digits. For example, the string “012345” could be entered in hex format as “30 31 32 33 34 35 00.” (Recall that the ASCII code for decimal digit x is $0x3x$, and that the end of a string is indicated by a null byte.)

The hex characters you pass to HEX2RAW should be separated by whitespace (blanks or newlines). We recommend separating different parts of your exploit string with newlines while you’re working on it. HEX2RAW supports C-style block comments, so you can mark off sections of your exploit string. For example:

```
48 c7 c1 f0 11 40 00 /* mov    $0x40011f0,%rcx */
```

Be sure to leave space around both the starting and ending comment strings (“/*”, “*/”), so that the comments will be properly ignored.

If you generate a hex-formatted exploit string in the file `exploit.txt`, you can apply the raw string to CTARGET or RTARGET in several different ways:

1. You can set up a series of pipes to pass the string through HEX2RAW.

```
unix> cat exploit.txt | ./hex2raw | ./ctarget
```

2. You can store the raw string in a file and use I/O redirection:

```
unix> ./hex2raw < exploit.txt > exploit-raw.txt
unix> ./ctarget < exploit-raw.txt
```

This approach can also be used when running from within GDB:

```
unix> gdb ctarget
(gdb) run < exploit-raw.txt
```

3. You can store the raw string in a file and provide the file name as a command-line argument:

```
unix> ./hex2raw < exploit.txt > exploit-raw.txt
unix> ./ctarget -i exploit-raw.txt
```

This approach also can be used when running from within GDB.

B Generating Byte Codes

Using GCC as an assembler and OBJDUMP as a disassembler makes it convenient to generate the byte codes for instruction sequences. For example, suppose you write a file `example.s` containing the following assembly code:

```
1 # Example of hand-generated assembly code
2     pushq    $0xabcdef          # Push value onto stack
3     addq     $17,%rax           # Add 17 to %rax
4     movl     %eax,%edx          # Copy lower 32 bits to %edx
```

The code can contain a mixture of instructions and data. Anything to the right of a '#' character is a comment.

You can now assemble and disassemble this file:

```
unix> gcc -c example.s
unix> objdump -d example.o > example.d
```

The generated file `example.d` contains the following:

```
1 example.o:          file format elf64-x86-64
2
3 Disassembly of section .text:
4
5 0000000000000000 <.text>:
6   0: 68 ef cd ab 00      pushq  $0xabcdef
7   5: 48 83 c0 11         add    $0x11,%rax
8   9: 89 c2              mov    %eax,%edx
```

The lines at the bottom show the machine code generated from the assembly language instructions. Each line has a hexadecimal number on the left indicating the instruction's starting address (starting with 0), while the hex digits after the ':' character indicate the byte codes for the instruction. Thus, we can see that the instruction `push $0xABCDEF` has hex-formatted byte code `68 ef cd ab 00`.

From this file, you can get the byte sequence for the code:

```
1 68 ef cd ab 00 48 83 c0 11 89 c2
```

This string can then be passed through `HEX2RAW` to generate an input string for the target programs.. Alternatively, you can edit `example.d` to omit extraneous values and to contain C-style comments for readability, yielding:

```
1   68 ef cd ab 00    /* pushq  $0xabcdef */
2   48 83 c0 11       /* add    $0x11,%rax */
3   89 c2            /* mov    %eax,%edx */
```

This is also a valid input you can pass through `HEX2RAW` before sending to one of the target programs.